


# 网络编程技术


---

UDP网络编程（二）

清华大学网研院

张千里



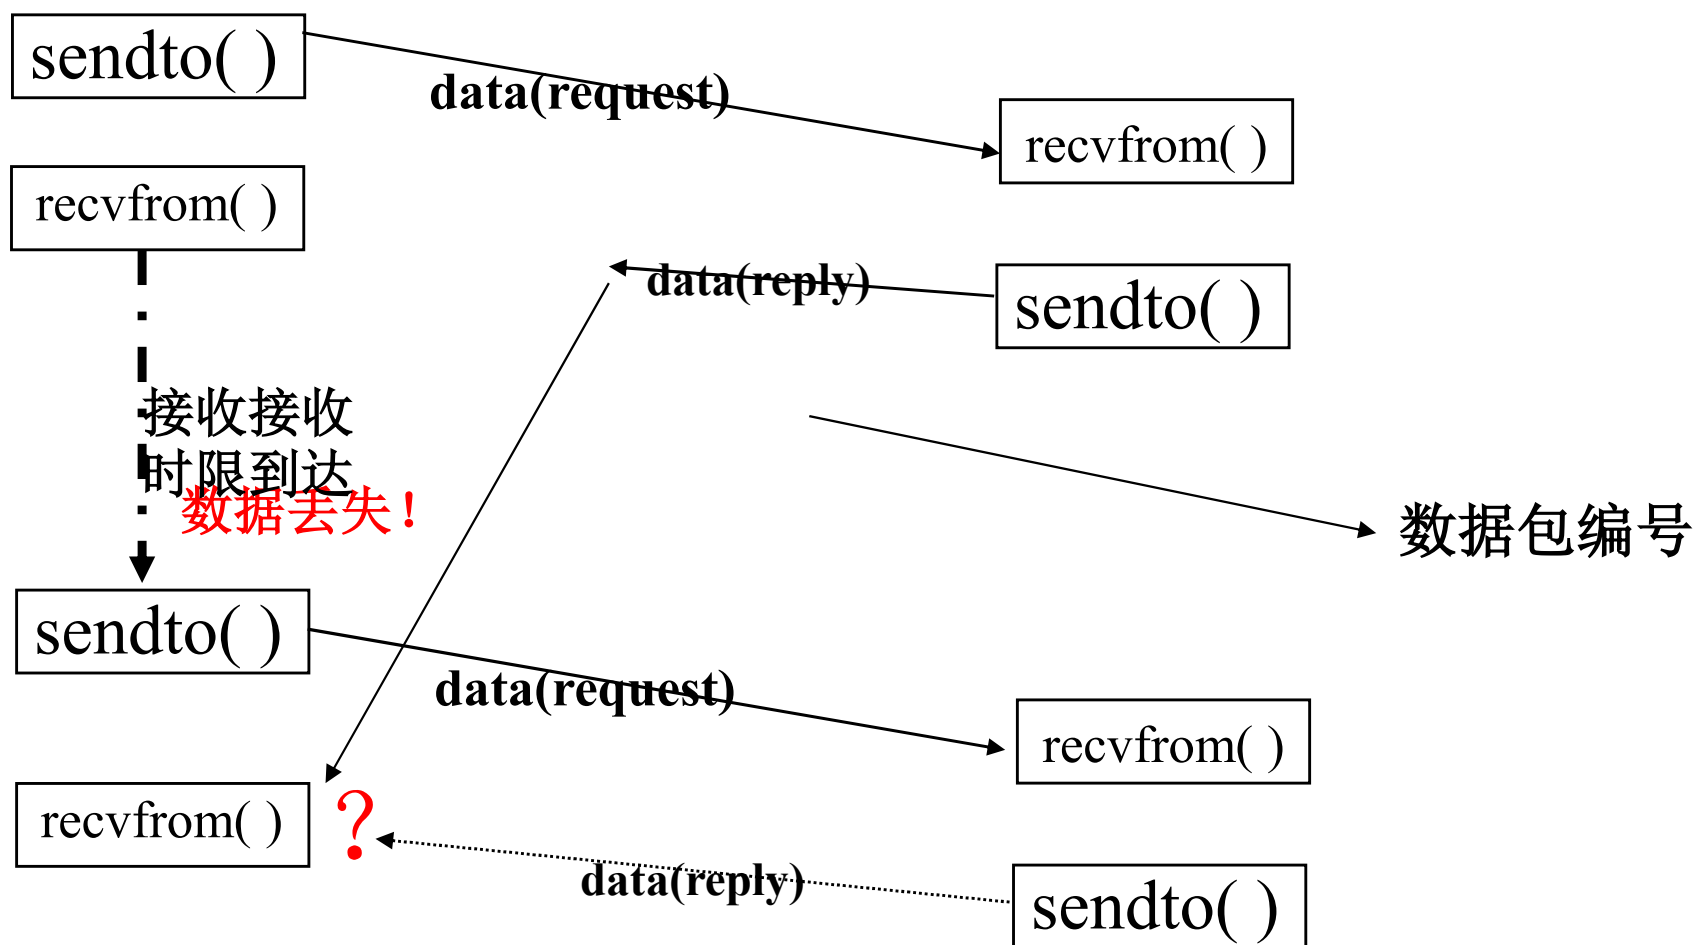




# 本节课内容

- 主要内容
  - **UDP套接字编程**
- 参考资料
  - **UNIX Network programming: 第8、21、22章**
    - Elementary UDP Sockets
    - Multicasting
    - Advanced UDP Sockets

# 可靠性举例（二）

## 数据对应





# 解决方法：使用标识号

- 使用标识号来进行请求/应答的对应，如**DNS**协议
  - **identification: 16 bit #**  
用于查询, 应答报文使用同样的 #
  - 而且，**DNS**响应中有问题项
- 当然也可以限定必须经过充分长时间后才能发下一个请求，但是这样会对应用产生较大限制

identification	flags
number of questions	number of answer RRs
number of authority RRs	number of additional RRs
questions (variable number of questions)	
answers (variable number of resource records)	
authority (variable number of resource records)	
additional information (variable number of resource records)	

↑  
12 bytes  
↓

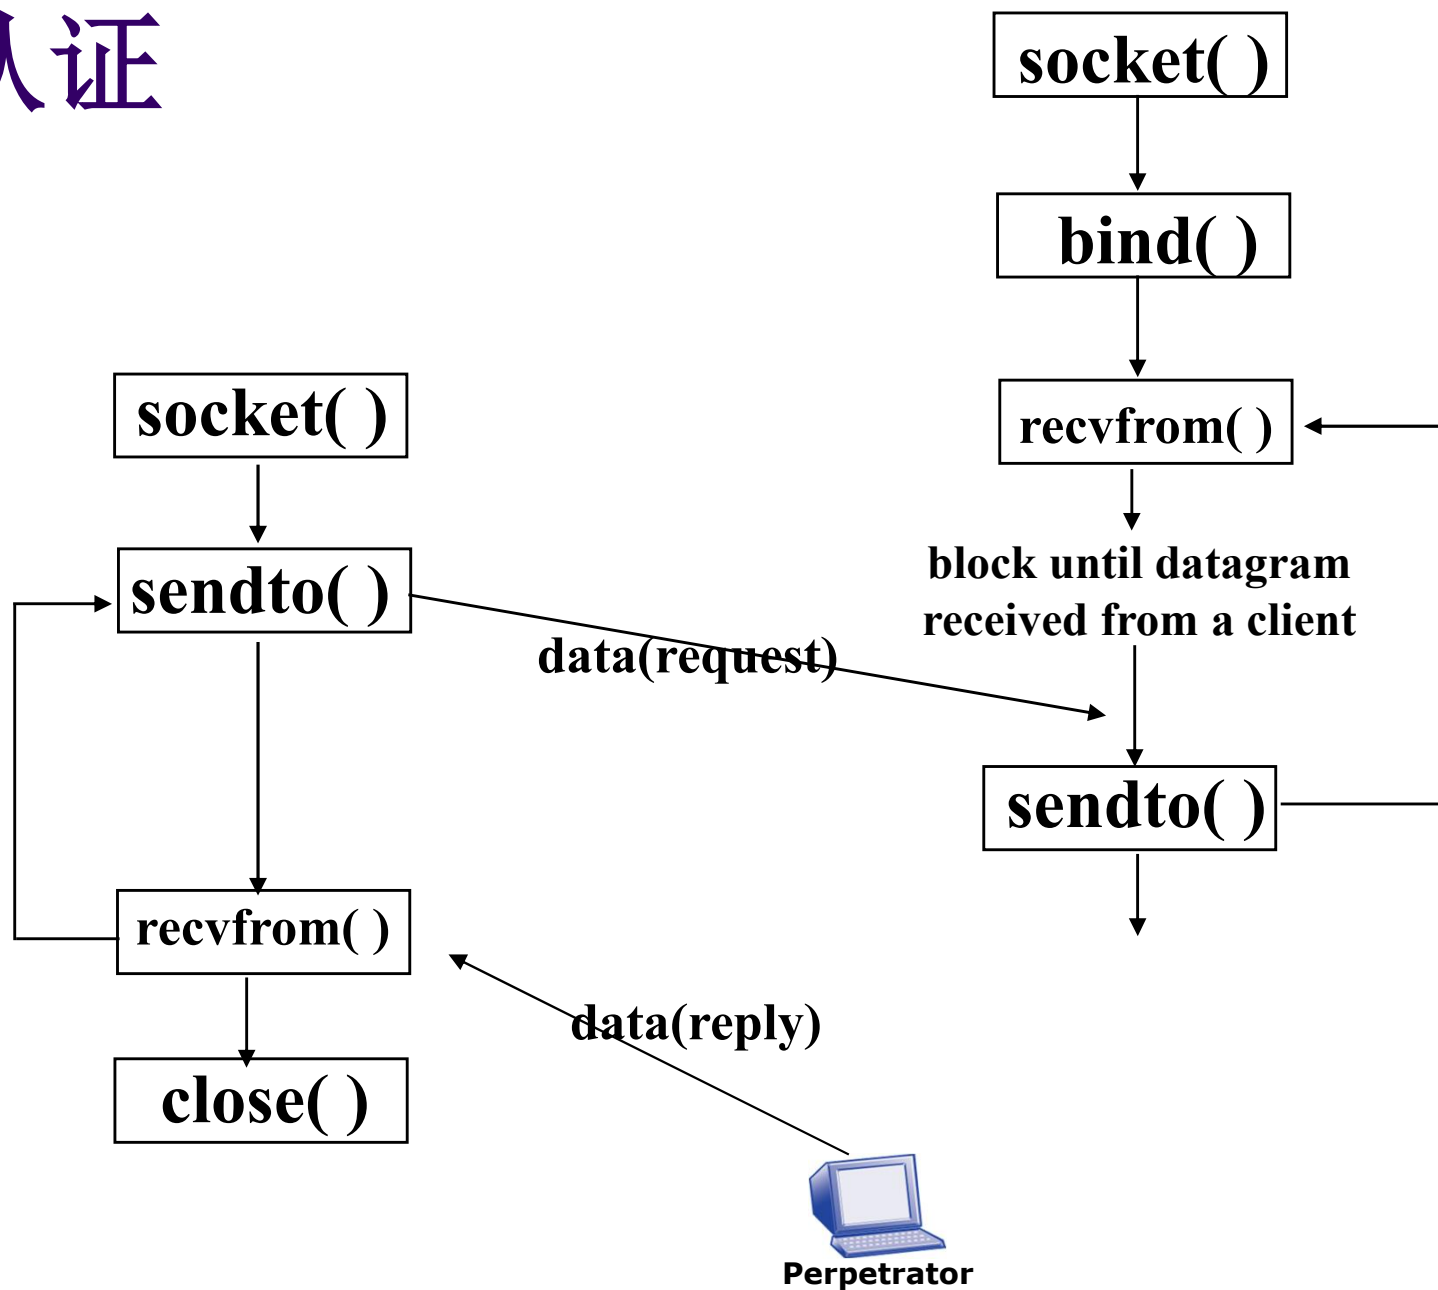



# 安全性

- 客户端所需要考虑的
  - 避免服务器假冒（源认证）
- 服务器所需要考虑的
  - 避免未认证的客户端的服务请求
  - 避免被拒绝服务攻击
  - 避免被利用进行拒绝服务攻击
  - 避免泄露客户的会话内容

# 安全性举例（一）

## 源认证






## lib/dg\_cli.c

```
#include "unp.h"
void dg_cli(FILE *fp, int sockfd, const SA *pservaddr, socklen_t servlen)
{
    int n;
    char sendline[MAXLINE], recvline[MAXLINE + 1];
    while (Fgets(sendline, MAXLINE, fp) != NULL) {
        Sendto(sockfd, sendline, strlen(sendline), 0, pservaddr, servlen);
        n = Recvfrom(sockfd, recvline, MAXLINE, 0, NULL, NULL);
        recvline[n] = 0; /* null terminate */
        Fputs(recvline, stdout);
    }
}

n = Recvfrom(sockfd, recvline, MAXLINE, 0, preply_addr, &len);
if (len != servlen || memcmp(pservaddr, preply_addr, len) != 0) {
    printf("reply from %s (ignored)\n", Sock_ntop(preply_addr, len));
    continue;
}
```

接到的数据不源于服务器！



上面的程序是否解决了源认证的问题？你有什么更好的建议？

正常使用主观题需2.0以上版本雨课堂

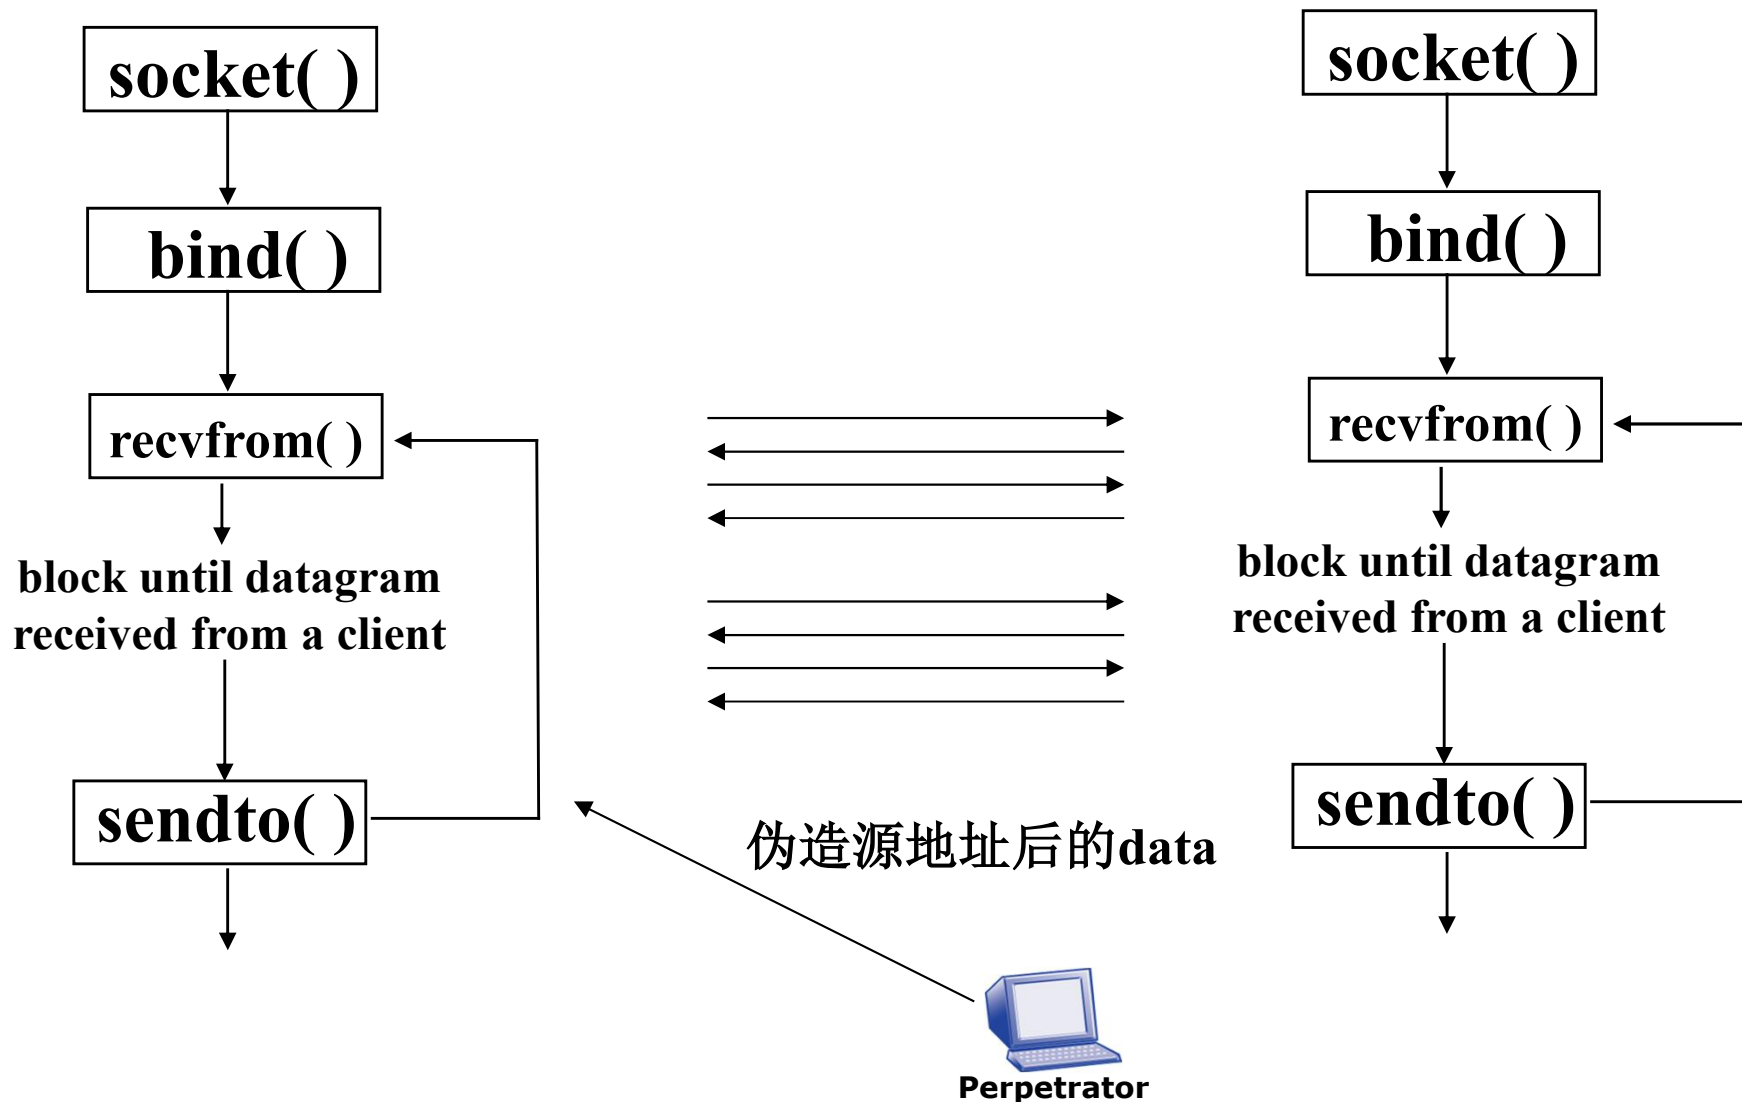

作答






# 安全性举例（二）

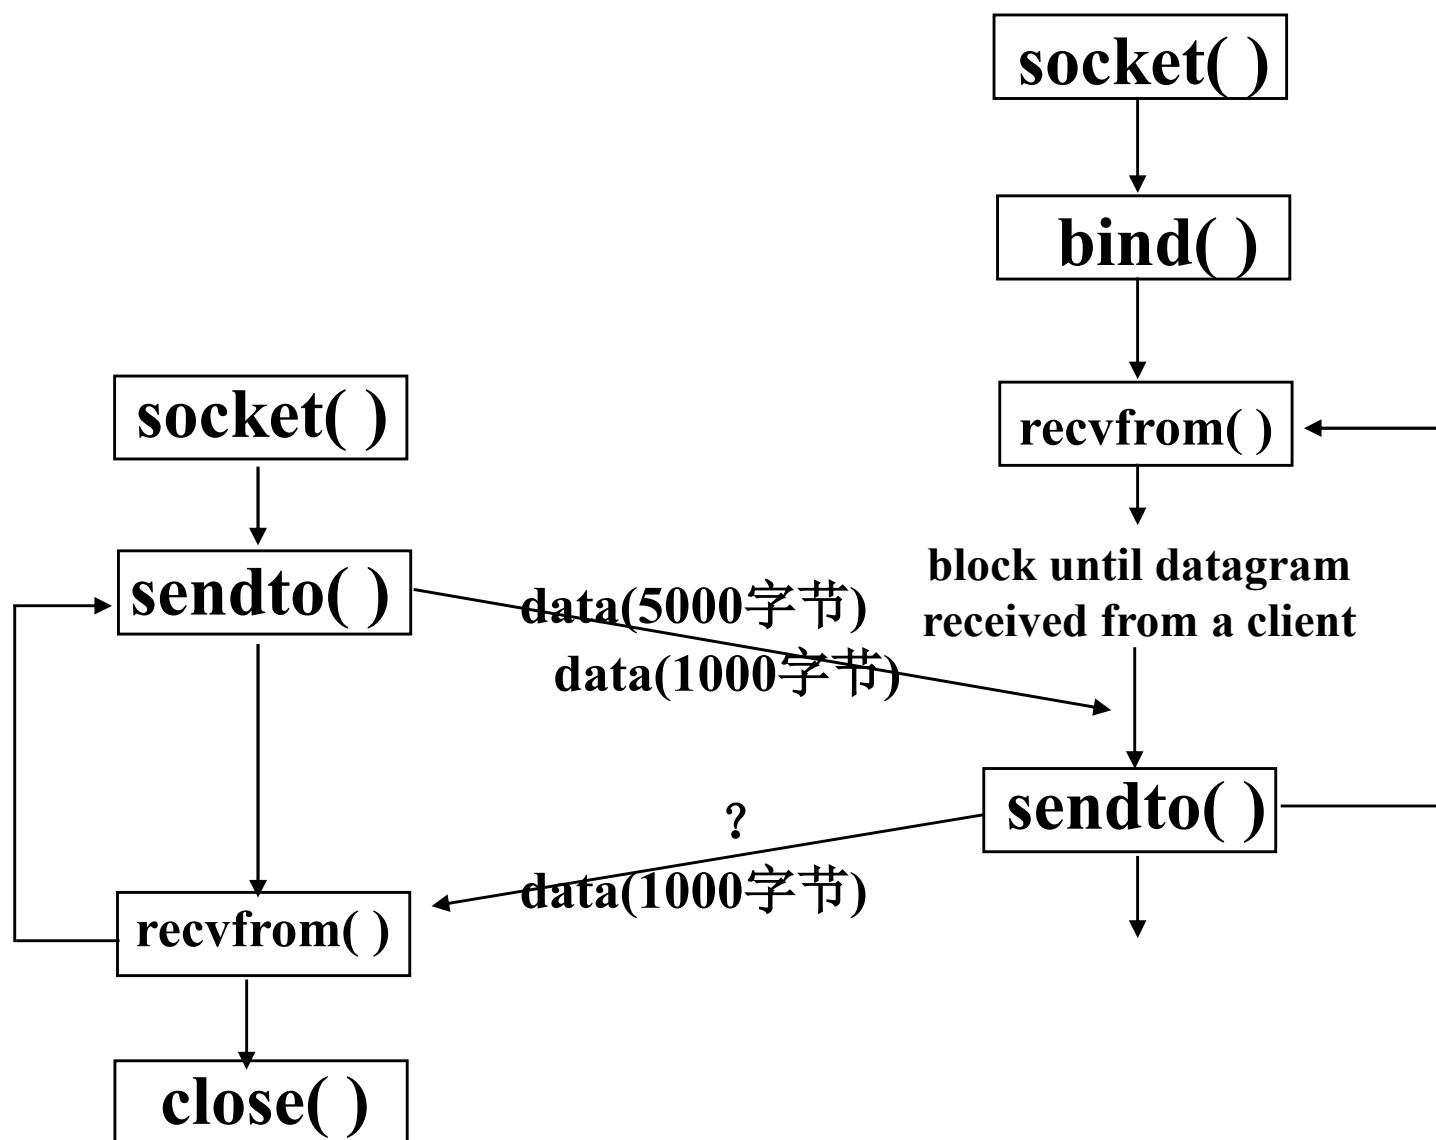
## 拒绝服务攻击






# 一致性

- 协议设计中的一致性考虑，举例：**ECHO**就是将用户输入的内容返回给用户
  - 问题1：协议如何进行切割？
    - 以回车符进行切割
    - 以一个报文进行切割
  - 问题2：长度是否有限制？
    - **MAXLINE**的影响
    - **UDP**最大报文长度的影响
    - **MTU**的影响
  - 问题3：何时返回？
    - 长请求需要接收到所有数据后返回
    - 每收到一部分后就立即返回
- 服务器实现应当严格遵循协议规范
  - 举例：错误处理，返回错误、直接忽略还是服务器退出？
  - 举例：超长报文，是否会返回部分？







# 报文接收recv等

```
int recv(int sd, void *buf, size_t len, int flags);  
int recvfrom(int sd, void *buf, size_t len, int flags, struct  
sockaddr *from, socklen_t *fromlen);
```

- **flags:** 功能选项，一般设置为0
  - **MSG\_TRUNC:** 指明数据报尾部数据已被丢弃，因为它比所提供的缓冲区需要更多的空间。
    - 一般而言，当收到比buf大的数据包时，剩余部分会被丢弃，返回的长度即buf的大小
    - 使用MSG\_TRUNC之后，返回的长度为整个数据包的长度


```
char buffer[MAX_SIZE];
struct sockaddr_in from;
struct sockaddr_in servaddr;
socklen_t fromlen;
int available_data;
int sock;
sock = socket(AF_INET, SOCK_DGRAM, 0);
fromlen = sizeof (from);
bzero(&servaddr, sizeof(servaddr));
servaddr.sin_family = AF_INET;
servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
servaddr.sin_port = htons(2000);
bind(sock, (struct sockaddr *) &servaddr, sizeof(servaddr));
available_data = recvfrom(sock, (char*)buffer, MAX_SIZE,
    MSG_TRUNC, (struct sockaddr *)&from, &fromlen);
if (available_data > MAX_SIZE) {
    fprintf(stderr, "UDP Packet is bigger than expected\n");
    exit(EXIT_FAILURE);
}
```





# 其他可选的标志


<i>flags</i>	Description	recv	send
MSG_DONTROUTE	不进行路由表查询		X
MSG_DONTWAIT	非阻塞	X	X
MSG_OOB	带外数据	X	X
MSG_PEEK	不清除接收队列	X	
MSG_WAITALL	阻塞到读满数据为止（除非信号或者其他数据到来）	X	




# 组播和广播

- 广播和组播不能在**TCP**下工作
  - 使用广播的协议
    - DHCP
    - NTP
  - IPv6大量使用组播

单播  
unicast




组播  
multicast



Type	IPv4	IPv6	TCP	UDP	# IP interfaces identified	# IP interfaces delivered to
Unicast	•	•	•	•	One	One
Anycast	*	•	Not yet	•	A set	One in set
Multicast	opt.	•		•	A set	All in set
Broadcast	•			•	All	All






# IPv4组播地址

- 预定义的组播地址
  - **224.0.0.1**: 子网里的所有系统
  - **224.0.0.2**: 子网里的所有路由器
  - ...


**Figure 21.3. Scope of IPv4 and IPv6 multicast addresses.**


Scope	IPv6 scope	IPv4	
		TTL scope	Administrative scope
Interface-local	1	0	
Link-local	2	1	224.0.0.0 to 224.0.0.255
Site-local	5	<32	239.255.0.0 to 239.255.255.255
Organization-local	8		239.192.0.0 to 239.195.255.255
Global	14	≤255	224.0.1.0 to 238.255.255.255



# IGMP

- 主机通过**IGMP**协议，向和自己相连的路由器报告组播组的加入、退出
- 当主机中的一个进程首次加入某个组时，主机向路由器发送报告，申请加入、退出组
- 路由器定期进行询问，主机响应
- 主机退出一个组时，不告知路由器，而是通过定期询问发现






# 组播套接字选项

Command	Datatype	Description
IP_MULTICAST_IF	struct in_addr	Specify default interface for outgoing multicasts
IP_MULTICAST_TTL	u_char	Specify TTL for outgoing multicasts
IP_MULTICAST_LOOP	u_char	Enable or disable loopback of outgoing multicasts
IPV6_MULTICAST_IF	u_int	Specify default interface for outgoing multicasts
IPV6_MULTICAST_HOPS	int	Specify hop limit for outgoing multicasts
IPV6_MULTICAST_LOOP	u_int	Enable or disable loopback of outgoing multicasts

Command	Datatype	Description
IP_ADD_MEMBERSHIP	struct ip_mreq	Join a multicast group
IP_DROP_MEMBERSHIP	struct ip_mreq	Leave a multicast group
IP_BLOCK_SOURCE	struct ip_mreq_source	Block a source from a joined group
IP_UNBLOCK_SOURCE	struct ip_mreq_source	Unblock a previously blocked source
IP_ADD_SOURCE_MEMBERSHIP	struct ip_mreq_source	Join a source-specific group
IP_DROP_SOURCE_MEMBERSHIP	struct ip_mreq_source	Leave a source-specific group
IPV6_JOIN_GROUP	struct ipv6_mreq	Join a multicast group
IPV6_LEAVE_GROUP	struct ipv6_mreq	Leave a multicast group
MCAST_JOIN_GROUP	struct group_req	Join a multicast group
MCAST_LEAVE_GROUP	struct group_req	Leave a multicast group
MCAST_BLOCK_SOURCE	struct group_source_req	Block a source from a joined group
MCAST_UNBLOCK_SOURCE	struct group_source_req	Unblock a previously blocked source
MCAST_JOIN_SOURCE_GROUP	struct group_source_req	Join a source-specific group
MCAST_LEAVE_SOURCE_GROUP	struct group_source_req	Leave a source-specific group




## 组播套接字选项（2）

- **IP\_ADD\_MEMBERSHIP, IPV6\_ADD\_MEMBERSHIP**
  - 通过指定的接口加入组
- **IP\_DROP\_MEMBERSHIP, IPV6\_DROP\_MEMBERSHIP**
  - 通过指定的接口退出组


```
Struct ip_mreq{  
    struct in_addr imr_multiaddr;  
    struct in_addr imr_interface;  
};
```

```
Struct ipv6_mreq{  
    struct in6_addr ipv6mr_multiaddr;  
    struct int ipv6mr_interface;  
};
```



## 组播套接字选项（3）

- **IP\_MULTICAST\_IF, IPV6\_MULTICAST\_IF**
  - 指定套接字发送组播数据包时使用的接口
  - 这个接口在IPv4中是in\_addr结构，而对于IPv6，则可以是接口序号（无符号整数）
- **IP\_MULTICAST\_TTL, IPV6\_MULTICAST\_TTL**
  - 设置组播数据包的TTL（IPv4）或者跳数上限（IPv6）
  - 未设置的话缺省为1
- **IP\_MULTICAST\_LOOP, IPV6\_MULTICAST\_LOOP**
  - 允许/禁止接受本地发送的组播报文（缺省允许）



# UNP中的例子 (lib/mcast\_join.c)

## mcast\_join及其他

```
#include "unp.h"
```

```
int mcast_join(int sockfd, const struct sockaddr *addr, socklen_t salen, const  
    char *ifname, u_int ifindex);
```

```
int mcast_leave(int sockfd, const struct sockaddr *addr, socklen_t salen);
```

```
int mcast_set_if(int sockfd, const char *ifname, u_int ifindex);
```

```
int mcast_set_loop(int sockfd, int flag);
```

```
int mcast_set_ttl(int sockfd, int ttl);
```

All above return :0 if ok, -1 on error

```
int mcast_get_if(int sockfd);
```


return : nonnegative interface index if OK, -1 error

```
int mcast_get_loop(int sockfd);
```

return : current loopback flag if OK, -1 error

```
int mcast_get_ttl(int sockfd);
```

return : current TTL or hop limit if OK, -1 error



```
int
mcast_join(int sockfd, const SA *grp, socklen_t grplen,
            const char *ifname, u_int ifindex)
{
#ifdef MCAST_JOIN_GROUP
    struct group_req req;
    if (ifindex > 0) {
        req.gr_interface = ifindex;
    } else if (ifname != NULL) {
        if ( (req.gr_interface = if_nametoindex(ifname)) == 0) {
            errno = ENXIO; /* i/f name not found */
            return(-1);
        }
    } else
        req.gr_interface = 0;
    if (grplen > sizeof(req.gr_group)) {
        errno = EINVAL;
        return -1;
    }
    memcpy(&req.gr_group, grp, grplen);
    return (setsockopt(sockfd, family_to_level(grp->sa_family),
                      MCAST_JOIN_GROUP, &req, sizeof(req)));
#endif
}
```



```
switch (grp->sa_family) {
case AF_INET: {
    struct ip_mreq      mreq;
    struct ifreq        ifreq;


    memcpy(&mreq.imr_multiaddr,
           &((const struct sockaddr_in *) grp)->sin_addr,
           sizeof(struct in_addr));

    if (ifindex > 0) {
        if (if_indextoname(ifindex, ifreq.ifr_name) == NULL) {
            errno = ENXIO; /* i/f index not found */
            return(-1);
        }
        goto doioclt;
    } else if (ifname != NULL) {
        strncpy(ifreq.ifr_name, ifname, IFNAMSIZ);

        if (ioctl(sockfd, SIOCGIFADDR, &ifreq) < 0)
            return(-1);
        memcpy(&mreq.imr_interface,
               &((struct sockaddr_in *) &ifreq.ifr_addr)->sin_addr,
               sizeof(struct in_addr));
    } else
        mreq.imr_interface.s_addr = htonl(INADDR_ANY);

    return(setsockopt(sockfd, IPPROTO_IP, IP_ADD_MEMBERSHIP,
                     &mreq, sizeof(mreq)));
}
```






```
case AF_INET6: {
    struct ipv6_mreq      mreq6;

    memcpy(&mreq6.ipv6mr_multiaddr,
           &((const struct sockaddr_in6 *) grp)->sin6_addr,
           sizeof(struct in6_addr));

    if (ifindex > 0) {
        mreq6.ipv6mr_interface = ifindex;
    } else if (ifname != NULL) {
        if ( (mreq6.ipv6mr_interface = if_nametoindex(ifname)) == 0) {
            errno = ENXIO; /* i/f name not found */
            return(-1);
        }
    } else
        mreq6.ipv6mr_interface = 0;

    return(setsockopt(sockfd, IPPROTO_IPV6, IPV6_JOIN_GROUP,
                     &mreq6, sizeof(mreq6)));
}
```




```
int
mcast_set_loop(int sockfd, int onoff)
{
    switch (sockfd_to_family(sockfd)) {
        case AF_INET: {
            u_char          flag;

            flag = onoff;
            return(setsockopt(sockfd, IPPROTO_IP, IP_MULTICAST_LOOP,
                               &flag, sizeof(flag)));
        }

#ifdef IPV6
        case AF_INET6: {
            u_int           flag;

            flag = onoff;
            return(setsockopt(sockfd, IPPROTO_IPV6, IPV6_MULTICAST_LOOP,
                               &flag, sizeof(flag)));
        }
#endif

        default:
            errno = EAFNOSUPPORT;
            return(-1);
    }
}
```



# mcast/main.c

```
if (argc != 3)
    err_quit("usage: sendrecv <IP-multicast-address> <port#>");

sendfd = Udp_client(argv[1], argv[2], (void **) &sasend, &salen);

recvfd = Socket(sasend->sa_family, SOCK_DGRAM, 0);

Setsockopt(recvfd, SOL_SOCKET, SO_REUSEADDR, &on, sizeof(on));

sarecv = Malloc(salen);
memcpy(sarecv, sasend, salen);
Bind(recvfd, sarecv, salen);


Mcast_join(recvfd, sasend, salen, NULL, 0);
Mcast_set_loop(sendfd, 0);

if (Fork() == 0)
    recv_all(recvfd, salen);                                /* child -> receives */


send_all(sendfd, sasend, salen);                            /* parent -> sends */
```


# UDP实际项目介绍

## n2n



- <https://www.ntop.org/products/n2n/>
- 设计需求
  - 边缘节点之间的流量加密
  - 边缘节点之间尽可能直接通讯，而不需要转发






# NAT

转发服务器

公网地址



Internet





NAT


公网地址



NAT

私有地址如  
192.168网段






# NAT穿越STUN

STUN回包告诉客户端  
公网IP和12345端口

私网

公网



What's  
my ip?

$x,y$



NAT

$[A,b]$

$A,b$

$S,t$




希望在5060端  
口接收数据

从5060端口发  
送请求STUN  
服务器

NAT映射端口为  
12345

现在数据可  
以直接发送



- ```

0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
! Version = 3      ! TTL              ! Flags                                !
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
4 ! Community ...
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
8 ! ... Community ...
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
12 ! ... Community ...
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
16 ! ... Community ...
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
20 ! ... Community
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+

```

```


0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-----+-----+-----+-----+-----+-----+-----+-----+
24 ! Source MAC Address
+-----+-----+-----+-----+-----+-----+-----+-----+
28 :                               ! Destination MAC Address
+-----+-----+-----+-----+-----+-----+-----+-----+
32 :
+-----+-----+-----+-----+-----+-----+-----+-----+
36 ! Socket Flags (v=IPv4)          ! Destination UDP Port
+-----+-----+-----+-----+-----+-----+-----+-----+
40 ! Destination IPv4 Address
+-----+-----+-----+-----+-----+-----+-----+-----+
44 ! Compress'n ID ! Transform ID ! Payload ...
+-----+-----+-----+-----+-----+-----+
48 !
+-----+-----+-----+-----+-----+-----+...

```




## n2n（三）协议状态

- （可选）EDGE向SN发送QUERY\_PEER
- EDGE向SN发送REGISTER\_SUPER
- 当发送数据时，向目标发送REGISTR，在成功后发送PACKET
- 如果打洞不成功，则PACKET还需要SN转发







请问以下的代码存在哪些问题？

```
receive = Recvfrom(server_socket, receive_buffer, SERVER_BUFFER, 0,  
                  (sockaddr *)&fromaddr, &addrlen);  
version = ntohl(*((unsigned int *)receive_buffer));  
switch(version){  
    case 1: {....};  
    case 2: {....};  
    default:  
        version = 1; // default to version 1  
        break;  
}  
...
```



# n2n (三) 协议实现

- 原语定义

- n2n\_typedefs.h

- 协议处理

- edge\_utils.c: process\_udp

- supernode.c: process\_udp

```
rc = select(max_sock + 1, &socket_mask, NULL, NULL, &wait_time);
```

```
now = time(NULL);
```

```
if(rc > 0) {
```

```
    // external udp
```

```
    if(FD_ISSET(sss->sock, &socket_mask)) {  
        struct sockaddr_in sender_sock;  
        socklen_t i;
```

```
        i = sizeof(sender_sock);
```

```
        bread = recvfrom(sss->sock, (void *)pktbuf, N2N_SN_PKTBUF_SIZE, 0 /*flags*/,  
                          (struct sockaddr *)&sender_sock, (socklen_t *)&i);
```

```
typedef struct n2n_PACKET {  
    n2n_mac_t      srcMac;  
    n2n_mac_t      dstMac;  
    n2n_sock_t     sock;  
    uint8_t        transform;  
    uint8_t        compression;  
} n2n_PACKET_t;
```

```
switch(msg_type) {  
    case MSG_TYPE_PACKET: {
```



# n2n (四) 协议解析

## ● wire.c

```
int decode_common (n2n_common_t * out,  
                  const uint8_t * base,  
                  size_t * rem,  
                  size_t * idx) {
```

```
    size_t idx0 = *idx;  
    uint8_t dummy = 0;
```

```
    decode_uint8(&dummy, base, rem, idx);
```

```
    if(N2N_PKT_VERSION != dummy) {  
        return -1;  
    }
```

```
    decode_uint8(&(out->ttl), base, rem, idx);  
    decode_uint16(&(out->flags), base, rem, idx);  
    out->pc = (out->flags & N2N_FLAGS_TYPE_MASK);  
    out->flags &= N2N_FLAGS_BITS_MASK;
```

```
    decode_buf(out->community, N2N_COMMUNITY_SIZE, base, rem, idx);
```

```
    return (*idx - idx0);
```

```
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1  
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+  
! Version = 3   ! TTL           ! Flags                               !  
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+  
4 ! Community ...                                     :  
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+  
8 ! ... Community ...                                 :  
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+  
12 ! ... Community ...                               :  
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+  
16 ! ... Community ...                               :  
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+  
20 ! ... Community                                   !  
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
```

```
        if(*rem < 2) {  
            return 0;  
        }
```

```
        *out  = ( base[*idx] & 0xff ) << 8;  
        *out |= ( base[1 + *idx] & 0xff );  
        *idx += 2;  
        *rem -= 2;
```

```
        return 2;
```






# 协议打包

```
*(base + (*idx))      = ( v >> 8) & 0xff  
*(base + (1 + *idx)) = ( v & 0xff );  
*idx += 2;
```

```
int encode_common (uint8_t * base,  
                  size_t * idx,  
                  const n2n_common_t * common) {  
  
    uint16_t flags = 0;  
  
    encode_uint8(base, idx, N2N_PKT_VERSION);  
    encode_uint8(base, idx, common->ttl);  
  
    flags = common->pc & N2N_FLAGS_TYPE_MASK;  
    flags |= common->flags & N2N_FLAGS_BITS_MASK;  
  
    encode_uint16(base, idx, flags);  
    encode_buf(base, idx, common->community, N2N_COMMUNITY_SIZE);  
  
    return -1;  
}
```

return 2;



Q&A